


GPU Programming in Computer Vision



Overview

- Project Overview
- ***Optical Flow***
 - Problem and Application
 - Hierarchical Horn and Schunck
 - Implementation on GPU
 - Results and Benefits
- Superresolution
 - Problem and Theory
 - Implementation on GPU
 - Results and Benefits
- Summary

Project Overview

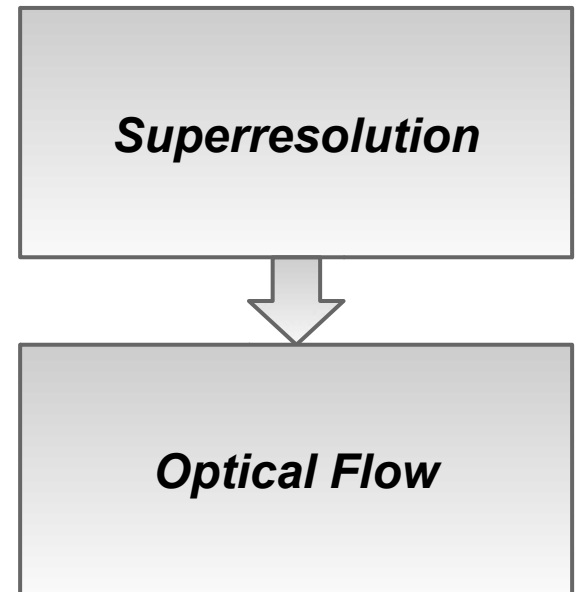
- ***Optical Flow***

- Understanding the theory
- Implementing a GPU Version of Hierarchical Horn and Schunck

- ***Superresolution***

- Understanding more theory ...
- Using results of optical flow

- Benchmarking different kinds of memory




Optical Flow – the Problem

"Optical flow or **optic flow** is the pattern of apparent motion of objects, surfaces, and edges in a visual scene caused by the relative motion between an observer (an eye or a camera) and the scene" (Wikipedia)

Applications:

- Motion Estimation
- Video Compression
- Time to Collision / Collision Avoidance
- Segmentation

Optical Flow



Horn and Schrunk Method

Global approach: *Horn und Schrunk Method*

Goal: Identification of a dense field of optical flux u for each pixel of given images I_1 and I_2 :

$$I_2(\mathbf{x} + u(\mathbf{x})) = I_1(\mathbf{x})$$

Together with the smoothness term we get the energy equation $E(u)$ to minimize:

$$E(u) = \int_{\Omega} (I_2(\mathbf{x} + u(\mathbf{x})) - I_1(\mathbf{x}))^2 + \lambda |\nabla u|^2 \quad d\mathbf{x}$$

with the gradient magnitude:

$$|\nabla u| = \sqrt{u_{1x}^2 + u_{1y}^2 + u_{2x}^2 + u_{2y}^2}$$

Horn and Schrunk Method

$E(u)$ is highly non-convex: do a Taylor expansion for the I_2 :

$$I_2(\mathbf{x} + u(\mathbf{x})) = I(\mathbf{x} + u(\mathbf{x}), t + 1) \approx I(\mathbf{x}, t) + \nabla I^\top u + \underbrace{I_t}_{\frac{dI}{dt}}$$

and plug it into the energy equation:

$$\begin{aligned} E(u) &= \int_{\Omega} (I(\mathbf{x}, t) + \nabla I^\top u + I_t - I(\mathbf{x}, t))^2 + \lambda |\nabla u|^2 \, d\mathbf{x} \\ &= \int_{\Omega} (\nabla I^\top u + I_t)^2 + \lambda |\nabla u|^2 \, d\mathbf{x} \end{aligned}$$

Horn and Schrunk Method

According to Euler-Lagrange calculus we get:

$$0 = I_x^2 u_1 + I_x I_y u_2 + I_x I_t - \lambda \Delta u_1$$

$$0 = I_x I_y u_1 + I_y^2 u_2 + I_y I_t - \lambda \Delta u_2$$

A robustification term is integrated into the equation:


$$E(u) = \int_{\Omega} \Phi_D((\nabla I^\top u + I_t)^2) + \lambda \Phi_R(|\nabla u|^2) \, d\mathbf{x}$$

This brought into LES and solved via SOR in a coarse to fine approach

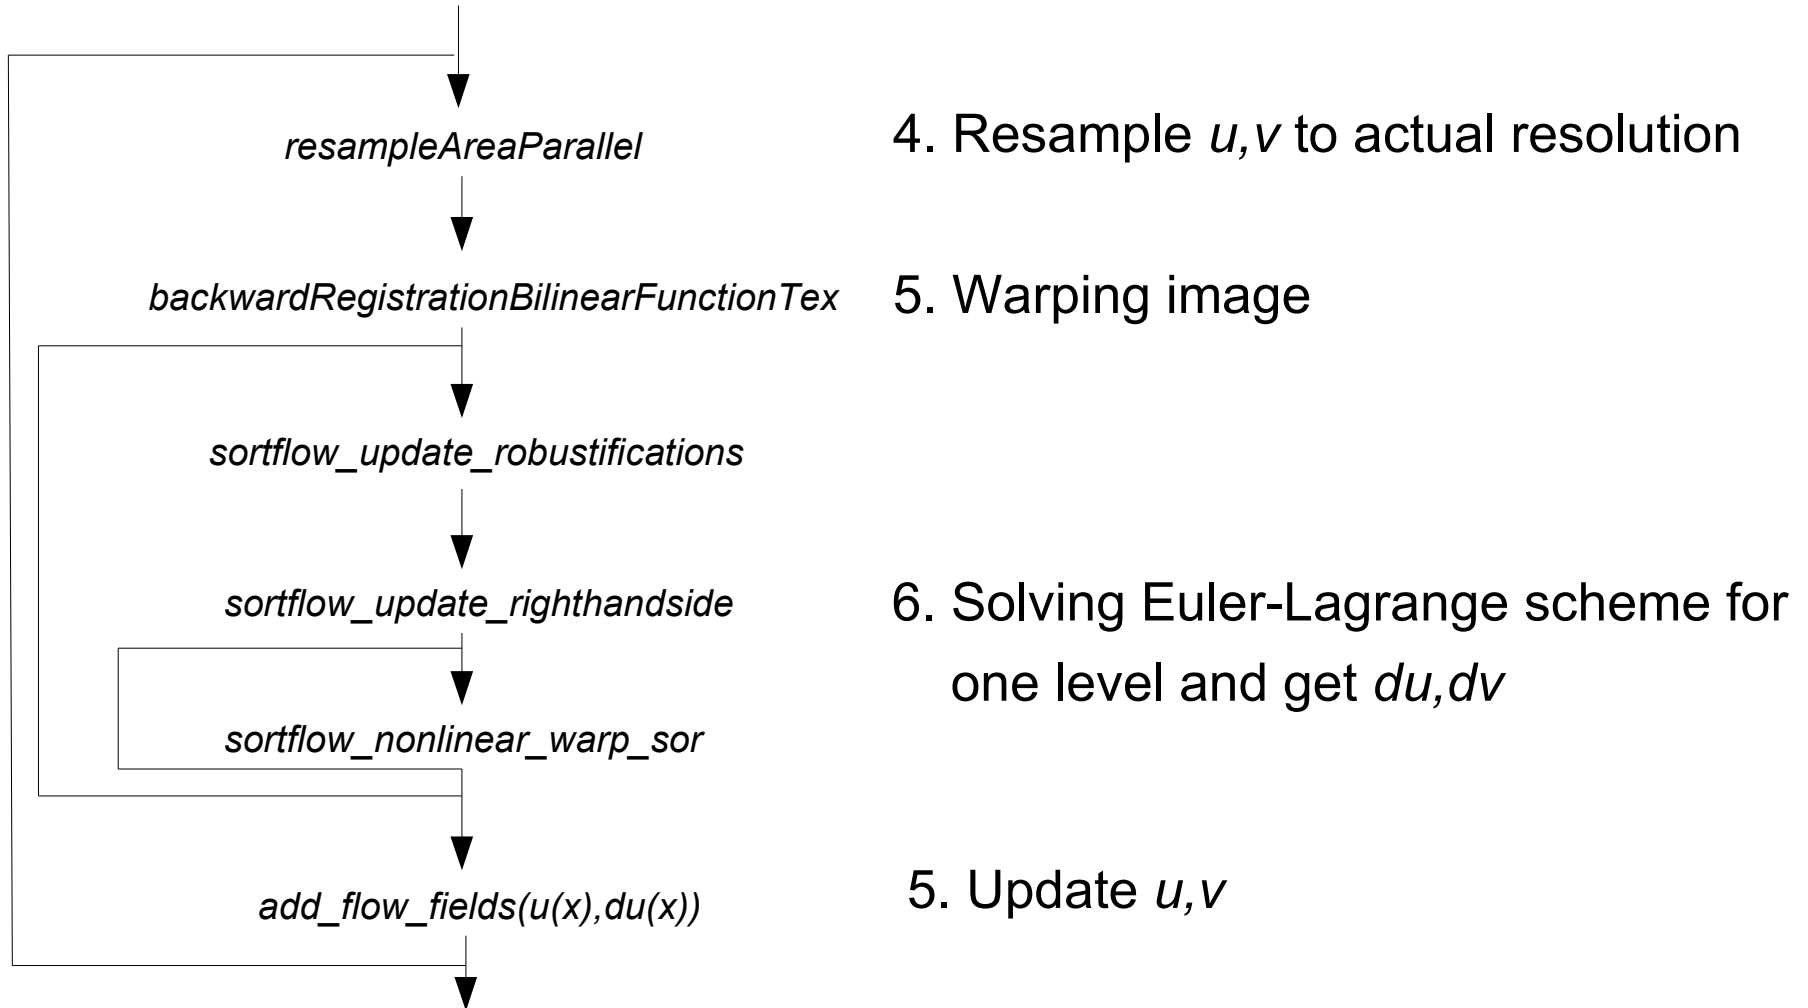
GPU Implementation

The (raw) algorithm:

1. Create the image pyramid (before)
2. Set u, v to zero
3. Start with the lowest resolution
4. Resample u, v to actual resolution
5. Warp Image with u, v
6. Solve Euler-Lagrange scheme for du, dv
7. Update u, v with du, dv
8. Go to higher resolution









GPU Implementation




Flow Results

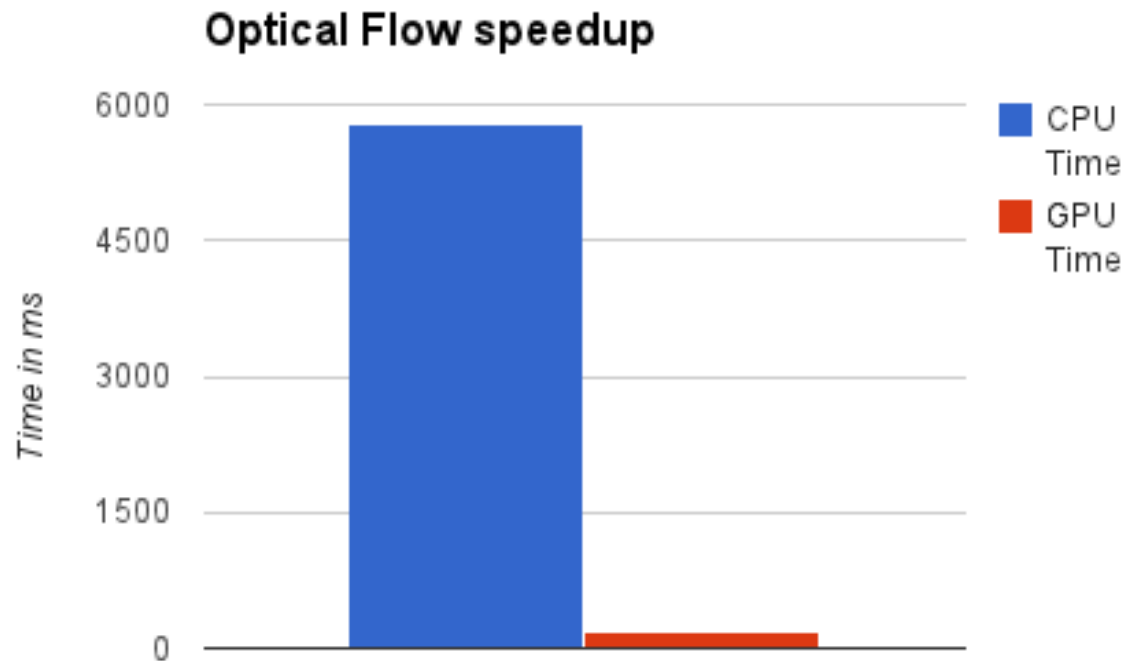
outer iterations



Where does the time go...



Results



average Speedup: 91,45

on a Tesla 1060 with 240 CUDA cores

Overview

- Project Overview
- Optical Flow
 - Problem and Application
 - Hierarchical Horn and Schunck
 - Implementation on GPU
 - Results and Benefits
- **Superresolution**
 - Problem and Theory
 - Implementation on GPU
 - Results and Benefits
- Summary

Superresolution


Superresolution is a term for a set of methods of upscaling video or images.

Idea: Generally, information is extracted from several low-res images and combined in a upscaled image.

Drawbacks:

- *Stills - if the object doesn't move, nothing can be extracted*
- *Very fast movements - blur problems*

Superresolution



Superresolution

Image I is degraded by a linear operator F .

Several degraded images, transformed by another linear operator T :

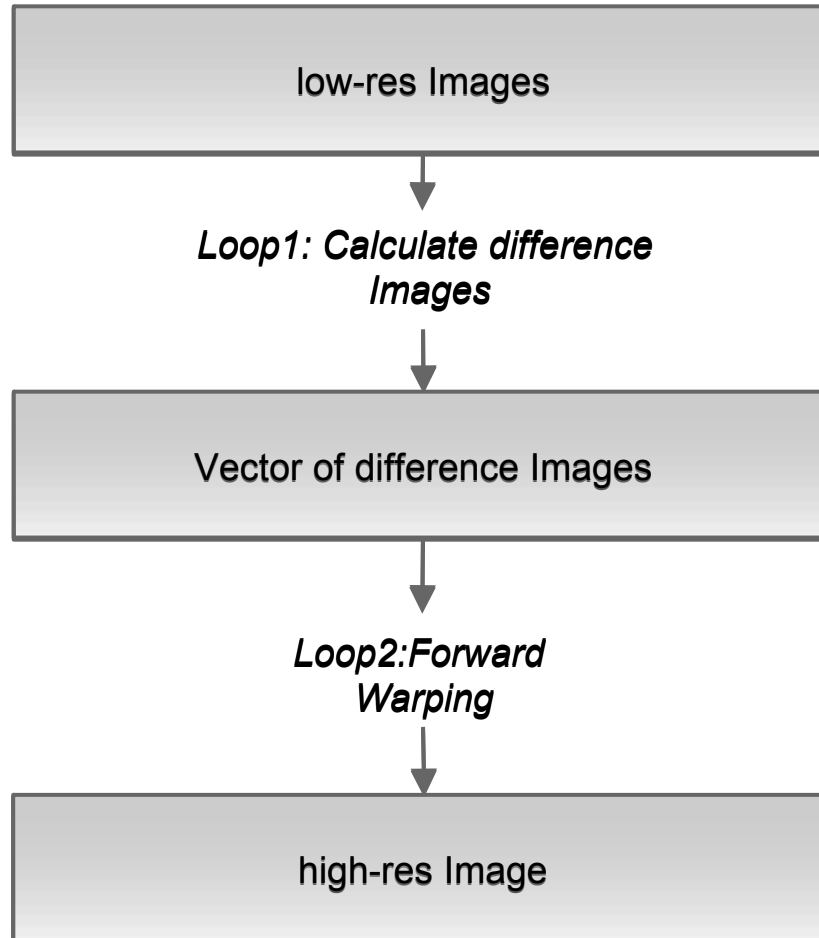
$$\begin{aligned} I_F^1 &= FT^1 I \\ I_F^2 &= FT^2 I \\ &\vdots \\ I_F^n &= FT^n I \end{aligned}$$

This is a large system of linear equations...

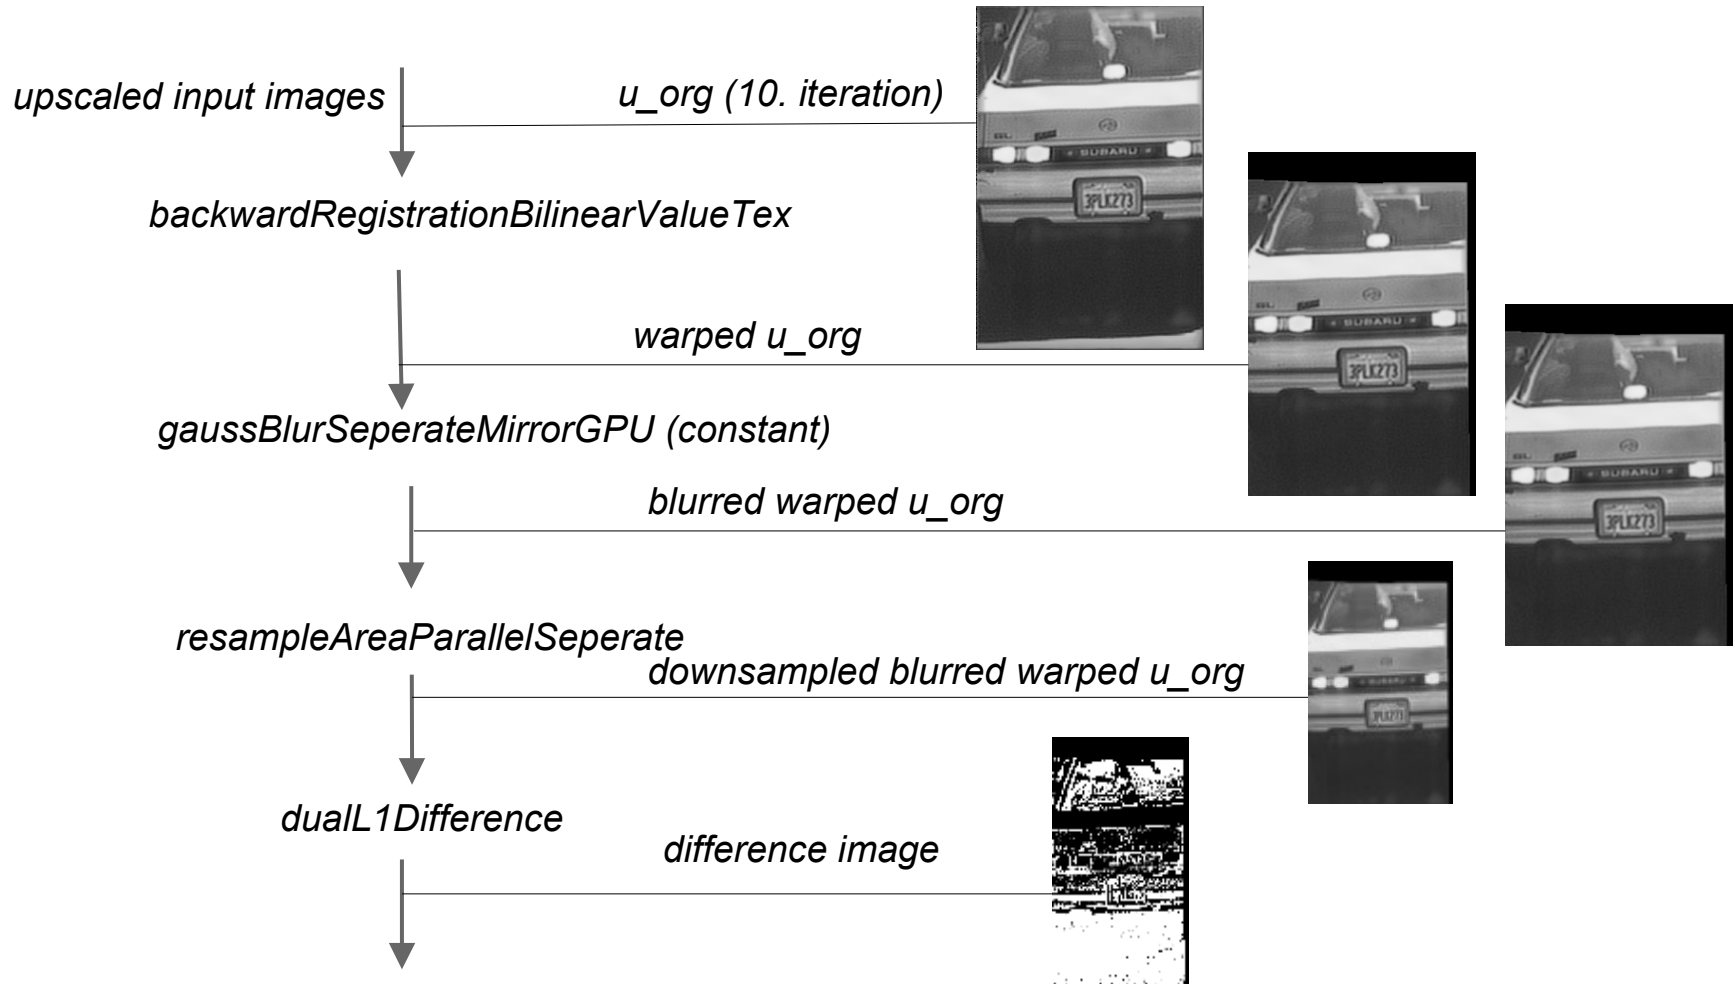
Finally, a energy function is fomulated:

$$E(I) = \sum_{i=1}^n \mu \|FT^i I - I_F^i\|_1 + \lambda \|\nabla I\|_1$$

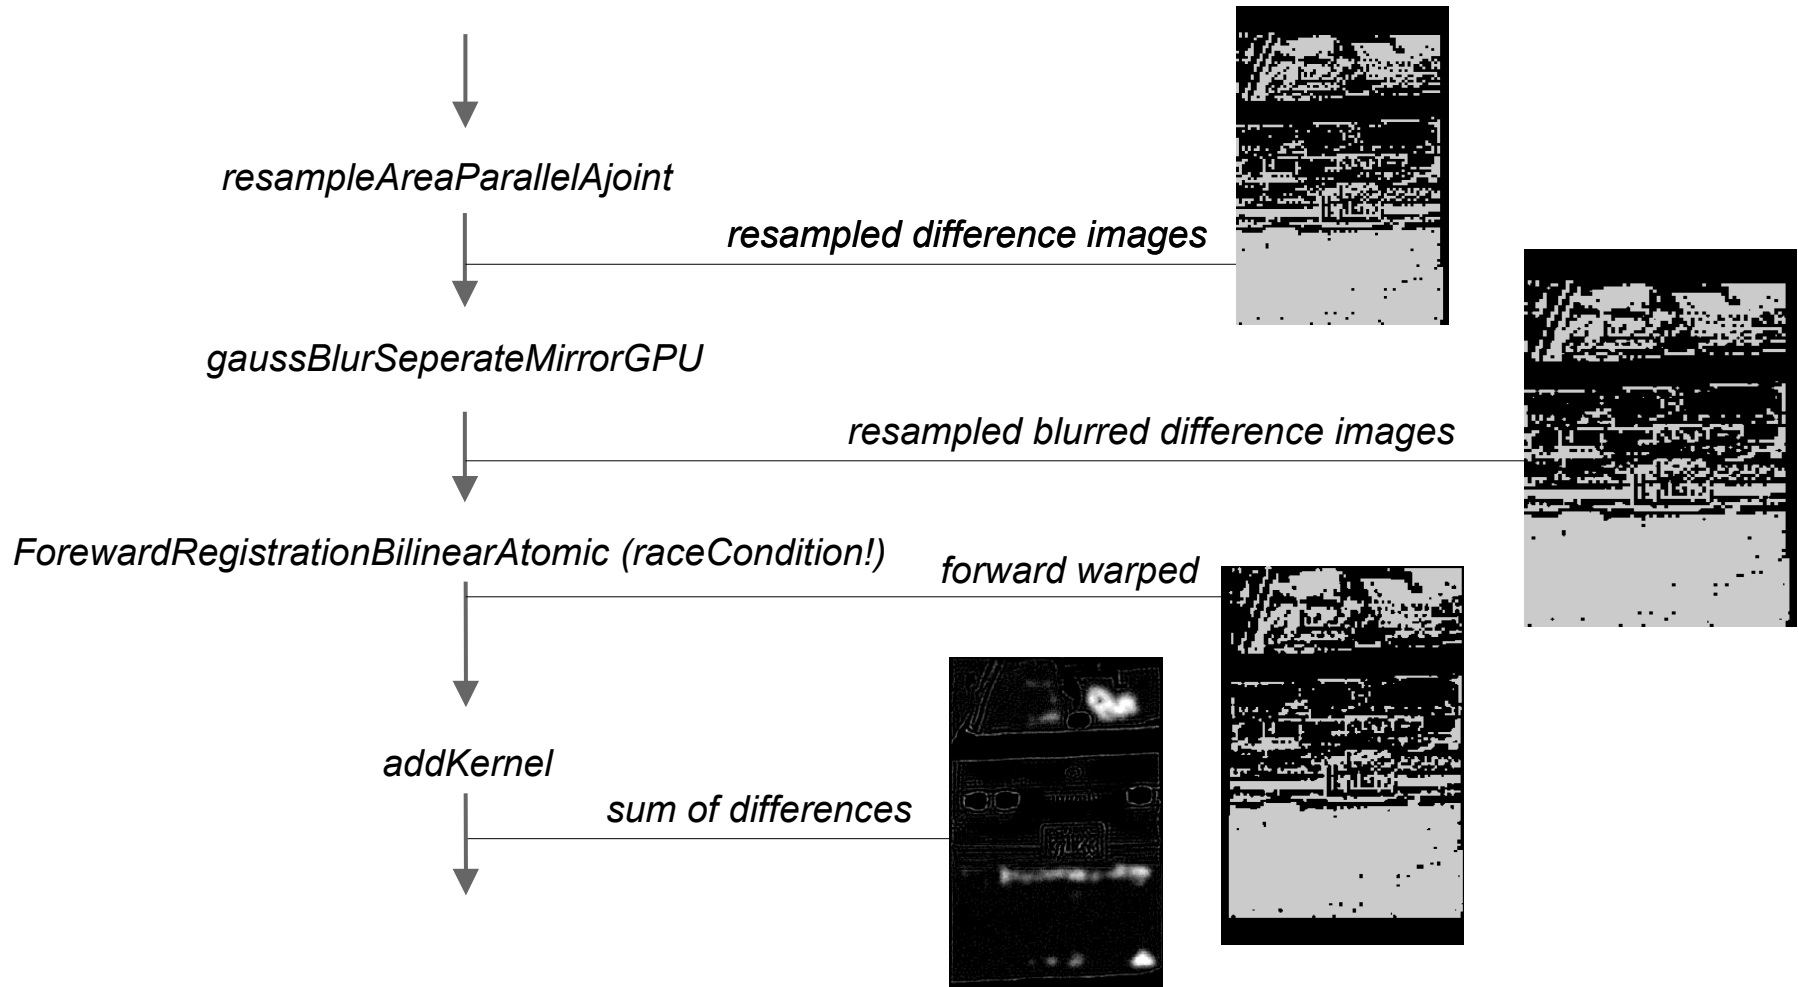
Superresolution



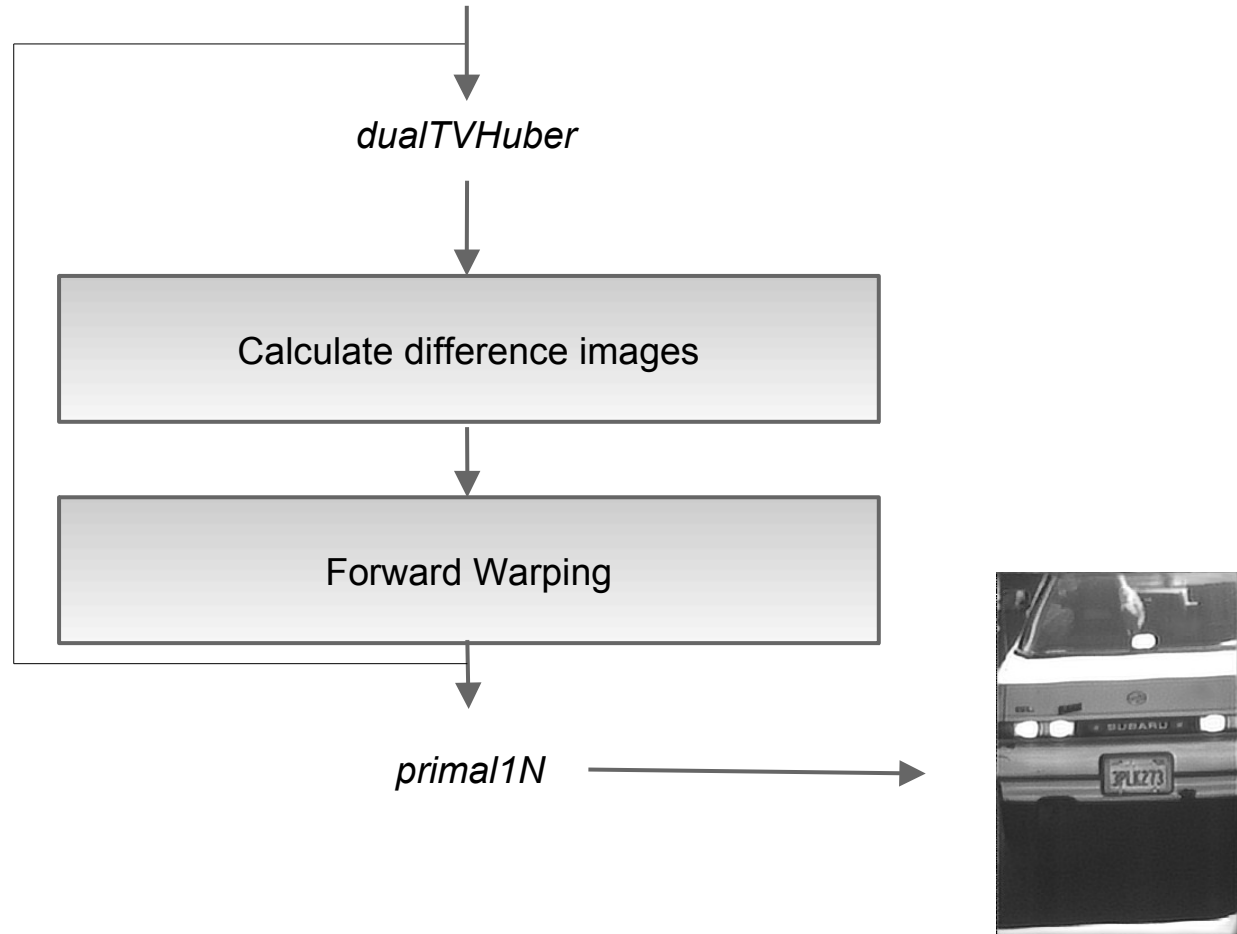
Superresolution: Loop 1




Superresolution: Loop 2




Superresolution



Where does the time go...



Results



average Speedup: 7,45

on a Tesla 1060 with 240 CUDA cores

Memory Experiments

- ***GaussBlur***
 - constant memory (kernel)
 - shared
 - global → **best**
 - texture → same as global
- ***BackwardWarping***
 - texture → **best**
 - global
- ***ForwardWarping***
 - atomics (race conditons)

Difficulties

- Understanding the algorithms
- Getting local environment run
- Understanding the framework
- Getting debugging possibilities
- How to get intermediate output
- Shared memory failure
- Incorrect working of CUDA printf()

Difficulties

```
float hx = (float)nx_in / (float)nx_out;
float factor = (float)(nx_out)/(float)(nx_in);

resampleAreaParallelSeparate_x<<< dimGrid, dimBlock >>>( in_g, help_g, nx_out, ny_in,
                                                         hx, pitchf1_in, pitchf1_out, factor);

// this cost us a lot of time -> resize grid to y_out
gridsize_y = (ny_out % LO_BH) ? ((ny_out / LO_BH)+1) : (ny_out / LO_BH);
dimGrid = dim3( gridsize_x, gridsize_y );

float hy = (float)ny_in / (float)ny_out;
factor = scalefactor*(float)ny_out / (float)ny_in;

resampleAreaParallelSeparate_y<<< dimGrid, dimBlock >>>( help_g, out_g, nx_out, ny_out,
                                                         hy, pitchf1_out, factor );
```

Lessons Learned

- Shared memory doesn't always work faster. It depends strictly upon amount of computation you do
- Use of different kind of memories (constant, texture) helps !
- As known, the speedup in the application comes with the cost of extra development efforts

Literature

Markus Unger , Thomas Pock , Manuel Werlberger , Horst Bischof, A convex approach for variational super-resolution, Proceedings of the 32nd DAGM conference on Pattern recognition, September 22-24, 2010, Darmstadt, Germany

Nils Papenberg , Andrés Bruhn , Thomas Brox , Stephan Didas , Joachim Weickert, Highly Accurate Optic Flow Computation with Theoretically Justified Warping, International Journal of Computer Vision, v.67 n.2, p.141-158, April 2006